

```
In [1]: # We will be using stepwise subset to model our EDA
# First we will check our trimmed data for any null values
import pandas as pd
df = pd.read_csv("retail_inventory_2023_2024Trimmed.csv")
```

```
In [2]: df.isna().sum()
```

```
Out[2]: Date                0
Store ID                  0
Product ID                0
Category                  0
Region                    0
Inventory Level            0
Units Sold                 0
Units Ordered              0
Demand Forecast            0
Price                      0
Discount                   0
Weather Condition          0
Holiday/Promotion          0
Competitor Pricing         0
Seasonality                0
Unnamed: 15                36600
Unnamed: 16                36600
dtype: int64
```

```
In [ ]: # We have two columns that have no values in it and that are also unnamed.
# Something in our preprocessing step messed up that format.
# We'll go ahead and drop those two columns now.
```

```
In [5]: df = df.iloc[:, :-2]
```


```
In [6]: df.isna().sum()
```

```
Out[6]: Date                0
Store ID                  0
Product ID                0
Category                  0
Region                    0
Inventory Level            0
Units Sold                 0
Units Ordered              0
Demand Forecast            0
Price                      0
Discount                   0
Weather Condition          0
Holiday/Promotion          0
Competitor Pricing         0
Seasonality                0
dtype: int64
```

```
In [ ]: # Our dependent variable is Units Sold. We'll focus on that and observe
# EDA through bar graphs, boxplots, and scatterplots
```


```
In [ ]: # We are going to see if the region, weather condition, and seasonality  
# have a part to play with units sold.
```

```
In [36]: # Region  
plt.figure(figsize=(8,5))  
sns.barplot(x='Region', y='Units Sold', data=df, estimator='mean')  
plt.title('Average Units Sold by Region')  
plt.show()
```




```
In [ ]: # Weather Condition
```

```
In [37]: plt.figure(figsize=(8,5))  
sns.barplot(x='Weather Condition', y='Units Sold', data=df, estimator='mean')  
plt.title('Average Units Sold by Weather Condition')  
plt.show()
```



```
In [38]: # Seasonality
plt.figure(figsize=(8,5))
sns.barplot(x='Seasonality', y='Units Sold', data=df, estimator='mean')
plt.title('Average Units Sold by Seasonality')
plt.show()
```



```
In [ ]: # The region, weather, and season don't have an effect on units sold.
# Regardless of these conditions, customers will still
# go into the store and purchase items.
```

```
In [7]: import matplotlib.pyplot as plt
import seaborn as sns

print(df.head())
print(df.info())
```

	Date	Store ID	Product ID	Category	Region	Inventory Level	\
0	2023-01-01	S001	P0001	Clothing	South	425	
1	2023-01-01	S001	P0002	Toys	North	240	
2	2023-01-01	S001	P0003	Clothing	North	330	
3	2023-01-01	S001	P0004	Clothing	North	383	
4	2023-01-01	S001	P0005	Toys	North	310	

	Units Sold	Units Ordered	Demand Forecast	Price	Discount	\
0	182	72	194.79	48.37	10	
1	91	135	99.06	69.97	20	
2	297	58	288.23	38.50	10	
3	10	166	26.69	31.03	10	
4	70	177	78.32	15.11	0	

	Weather Condition	Holiday/Promotion	Competitor Pricing	Seasonality
0	Rainy	1	49.99	Autumn
1	Cloudy	0	66.91	Winter
2	Snowy	1	39.10	Summer
3	Sunny	1	32.01	Winter
4	Rainy	0	14.21	Winter

```
<class 'pandas.DataFrame'>
```

```
RangeIndex: 36600 entries, 0 to 36599
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	36600 non-null	str
1	Store ID	36600 non-null	str
2	Product ID	36600 non-null	str
3	Category	36600 non-null	str
4	Region	36600 non-null	str
5	Inventory Level	36600 non-null	int64
6	Units Sold	36600 non-null	int64
7	Units Ordered	36600 non-null	int64
8	Demand Forecast	36600 non-null	float64
9	Price	36600 non-null	float64
10	Discount	36600 non-null	int64
11	Weather Condition	36600 non-null	str
12	Holiday/Promotion	36600 non-null	int64
13	Competitor Pricing	36600 non-null	float64
14	Seasonality	36600 non-null	str

```
dtypes: float64(3), int64(5), str(7)
```

```
memory usage: 4.2 MB
```

```
None
```

```
In [ ]: # Here we can see we have a couple of numerical and categorical values in this
# data set Some categories are related to each other, like weather condition
# and Seasonality. Going to start showing some visuals to help us determine
# what might be important to keep for our multiple linear regression

# There are also 36,600 recorded units sold in the year of 2023-2024. There
# were five stores in total with recorded data. Each store that recorded data
# for a certain day had different departments ranging from toys to clothing
# to furniture. Typical departments you'd find in a store
```


```
In [ ]: # Now time to display some histograms to understand the distribution
# and skewness
```

```
In [18]: import matplotlib.pyplot as plt

numeric_cols = df.select_dtypes(include=['float64', 'int64']).columns
axes = df[numeric_cols].hist(figsize=(12,10), bins=15)
plt.subplots_adjust(hspace=0.4, wspace=0.3)

for ax in axes.flatten():
    ax.set_xlabel("Value")
    ax.set_ylabel("Frequency")
plt.suptitle("Histograms of Forecasted Retail Store Variables")
plt.show()
```

Histograms of Forecasted Retail Store Variables



```
In [9]: print(df[numeric_cols].skew())
```


```
Inventory Level      0.011593
Units Sold           0.900578
Units Ordered        -0.000461
Demand Forecast      0.891590
Price                0.002817
Discount             -0.002700
Holiday/Promotion    0.017378
Competitor Pricing   0.001928
dtype: float64
```

```
In [ ]: # The frequency for units sold, our dependent variable, has a decreased trend
# going down. Overall this indicates the supply and demand theory of higher
# supply, less demand. The demand forecast backs up this theory for the
# most part too with the overall decrease trend as demand reaches its
# threshold. Besides these two variables, we can see that the discount
# seems to be only in groups of 5% discounts. 5%, 10%, 15%, and 20%.
# The holiday/promotion seems to be hot encoded. It's either a
# holiday/promotion or not. Only units ordered and discount are left
# skewed based on their negative skewness. The rest are right skewed.
# Visually speaking, it seems competitor pricing is the closest to
```

```
# being a nice bell shaped curve. Making it the closest to normal
# distribution
```

```
In [ ]: # Boxplots to detect outliers and confirm skewness
```


```
In [21]: plt.figure(figsize=(12,10))
sns.boxplot(data=df[numeric_cols])
plt.title('Boxplots Forecasted Retail Store Variables')
plt.xticks(rotation=45)
plt.show()
```




```
In [ ]: # As we can see visually, units sold and demand forecast tend to have outliers
# while the rest don't. This is due to the nature of supply and demand. What
# can attest for these reasons is the goal of our MLR predictions. Demand
# spikes, out of stock inventory, and expired/long shelf life products cause
# these volatile issues. We wish to prevent stores from losing revenue or
# having too much inventory to prevent long shelf life, so the need for a
# good prediction is evident. It's also quite interesting to see how the
# majority of variables right skewed visually. Since they are
# gravitating towards the bottom of the quartiles.
```

```
In [ ]: # Now it's time to view the bar graphs via different stores. We want to  
# analyze the mean of units sold in each store to see how the stores  
# differ. We will also take into account how each product is sold as  
# well
```

```
In [22]: plt.figure(figsize=(8,5))  
sns.barplot(x='Store ID', y='Units Sold', data=df, estimator='mean')  
plt.title('Average Units Sold by Store')  
plt.show()
```




```
In [24]: plt.figure(figsize=(12,8))  
sns.barplot(x='Product ID', y='Units Sold', data=df, estimator='mean')  
plt.title('Average Units Sold by Store')  
plt.xticks(rotation=45)  
plt.show()
```





```
In [25]: # The average units sold at each store seem relatively equal. Along with the  
# products too! Our data has a roughly, good representation of each store and  
# product. To confirm, we will look at the store counts and products too.  
# Just to really show our data is well represented which means our model will  
# handle/predict new data accurately as well
```

```
In [26]: sns.countplot(x='Store ID', data=df)  
plt.title('Count of Observations by Store ID')  
plt.show()
```




```
In [28]: plt.figure(figsize=(12,8))
sns.countplot(x='Product ID', data=df)
plt.title('Count of Observations by Product ID')
plt.xticks(rotation=45)
plt.show()
```



```
In [ ]: # Finally for visuals we will be making a correlation heat map. This will  
# help us identify what factors have strong correlation to help us decided  
# what variables we should choose for our multiple linear regression model.  
# There doesn't seem to be a strong bias from either a unique store or product
```

```
In [30]: plt.figure(figsize=(12,8))  
sns.heatmap(df[numeric_cols].corr(), annot=True, cmap='RdYlGn')  
plt.title('Correlation Matrix of Forecasted Retail Store Variables')  
plt.show()
```



```
In [31]: # Our focus will be units sold. here there are a lot of red or "weak" correlations
# between the majority of these variables. While looking at units sold though, we
# can see the green/light green or "stronger" relationships with inventory level
# and demand forecast. With a slightly weaker relationship between units ordered
# and holiday/promotion. This is indicating that our main focuses are going to be
# supply vs demand. With the help of external factors like holiday/promotion,
# discount, and units ordered, our multiple linear regression model is going to
# take a look at crafting the best prediction outcome based on demand of the
# product. This will be able to help us predict how much stock
# to buy because we will be able to predict how many units are going to be sold.
```

```
In [ ]: # Finally we will run a describe analysis to view the overall numerical values
```

```
In [33]: print(df.describe())
```

	Inventory Level	Units Sold	Units Ordered	Demand Forecast \
count	36600.000000	36600.000000	36600.000000	36600.000000
mean	274.149945	136.173716	110.142678	141.208696
std	129.958751	108.784554	52.194365	109.081550
min	50.000000	0.000000	20.000000	-9.970000
25%	162.000000	48.000000	65.000000	53.487500
50%	273.000000	107.000000	110.000000	112.980000
75%	386.000000	203.000000	155.000000	207.977500
max	500.000000	499.000000	200.000000	518.550000

	Price	Discount	Holiday/Promotion	Competitor Pricing
count	36600.000000	36600.000000	36600.000000	36600.000000
mean	54.988310	10.003279	0.495656	55.012043
std	26.076935	7.102970	0.499988	26.259164
min	10.000000	0.000000	0.000000	5.340000
25%	32.397500	5.000000	0.000000	32.437500
50%	54.870000	10.000000	0.000000	54.920000
75%	77.742500	15.000000	1.000000	77.730000
max	99.990000	20.000000	1.000000	104.940000

```
In [34]: # These numbers prove what we saw visually in terms of outliers and the
# distribution of data. Based on the mean and medians, we can see that the
# units ordered and discount are slightly left skewed. While the other
# variables are slightly right skewed. Units sold and demand forecast are
# the most skewed out of all the variables. Indicating most of our data is
# somewhat normally distributed. The standard deviation for our numerical
# values are higher for demand forecast and units sold versus the other
# variables. This explains the heavy skew that associates
# with demand forecast and units sold
```

```
In [35]: # Final thoughts
```

```
In [ ]: # Our data is well represented across all five stores and the total products that
# are being sold. Taking into consideration our high skewness with demand forecast,
# we might need to log transform our data. We are going to recommend not to use
# demand forecast. It is too highly correlated with units sold and could cause
# multicollinearity. On top of that, we would consider including inventory level,
# as well as holiday/promotion, discount, and units ordered. Given their slightly
# better correlation compared to the other variables.
# Finally the categorical variables, region, weather conditions, and seasonality,
# don't seem relevant. Demand of these products at all five stores still holds
# true regardless of region, weather, or season. There is a slight favorability
# towards the east region, sunny day, and autumn weather, but it's not enough
# to have an impact on units sold.
```